

Cloudifier Virtual Apps

Virtual desktop predictive analytics apps environment based on GPU computing framework

Andrei Ionut DAMIAN (*Author*)

Cloudifier SRL
Bucharest, Romania
damian@cloudifier.net

Alexandru PURDILA (*Author*)

Cloudifier SRL
Bucharest, Romania
alex@cloudifier.net

Nicolae TAPUS (*Author*)

University “Politehnica” Bucharest
Bucharest, Romania
ntapus@cs.pub.ro

Abstract — The need for systems capable of conducting inferential analysis and predictive analytics is ubiquitous in a global information society. With the recent advances in the areas of predictive machine learning models and massive parallel computing a new set of resources is now potentially available for the computer science community in order to research and develop new truly intelligent and innovative applications. In the current paper we present the principles, architecture and current experimentation results for an online platform capable of both hosting and generating *intelligent* applications.

Keywords— *artificial intelligence, machine learning, virtual desktop, predictive analytics*

- users in need of a online application marketplace fully oriented on predictive analytics applications;
- software developers in need of a legacy apps automatic software translation tools including machine learning templates;
- scientific users both from academic and commercial environment looking for an online application framework and execution environment in order to conduct predictive experiments in areas such as healthcare, finance and environment;

I. INTRODUCTION

The entire field of Artificial Intelligence tends to become more and more about improving predictions and predictive models as Goldfarb et al [1] argue in the paper “Managing the Machines”. Predictive modelling can now be used for extremely varied tasks ranging from constructing expert system chat-bots [2] based on Long Short Term Memory deep recurrent neural networks that would use natural language processing capability (reading and writing) up to melanoma recognition in dermoscopy images [3] based on deep convolutional networks.

In our paper, we will propose several approaches: (i) a web-services based model for the deployment of model training and prediction on GPU massive parallel computing backend and (ii) employment of automatic software design, development and deployment for user experience modules based on a pipeline shallow machine learning model [4] and (iii) a unified architecture and underlying framework for generation and hosting predictive analytics applications.

Within our paper we will argue that our proposed research and experimentation will fuel the development and implementation of a production-grade system capable of addressing in the same time the needs of several categories of customers such as:

II. RELATED WORK

Our work relates to the most recent advances both in the field of High Performance Computing and Machine Learning, particularly in the areas of machine learning model deployment on massive parallel processing infrastructures based on numerical computing cores.

A. GPU based High Performance Computing

The usage of GPU based computing is currently mostly applied in deep learning – both research experiments and production development – however little to no attention is given to the employment of GPU parallel computing resources in shallow machine learning techniques such as decision trees, multi-variate linear and logistic regression, Naive-Bayes models. Nonetheless the current advances and the state-of-the-art prediction scores achieved by the deep learning scientific experiments give little room to further scientific development and advances of known shallow models techniques. Even more, companies such as Microsoft, Google or IBM currently propose online framework-engines that allow on-demand web-based prediction model creation and experimentation [5] that are entirely used for off-line training and prediction experiments, running on CPU based infrastructures. Taking these presented aspects into account we are proposing a set of new modern approaches that will potentially empower shallow machine

learning models with the capability using GPU based parallel computing resources. Also, another of the supporting aspects that we will use in our approach is based on the most important features of shallow machine learning models: the ability of the model to self-explain itself (albeit at the expense of prediction power shown by deep models with massive hidden and less than explanatory weight spaces).

B. Architectural roots

The proposed architecture is based on our previous work conducted in the area of automatic software application migration and machine-learning assisted automatic programming [4]. Based on our automatic software migration architecture [4] combined with the latest advances in large-scale machine learning on heterogeneous distributed systems [6] we will be able to achieve the following desiderates during the production phase of our project:

- Create a scalable architecture for a user-friendly online environment where external users will be able to quickly generate and customize application starting from existing legacy apps;
- Combine the power of the automatic application generation features with template-based models in order to create a business horizontal application repository that will include apps such as (a) multi-vertical customer churn prediction models; (b) multi-vertical customer behavior inference model; (c) business flow generated data exploration stub apps; (d) machine learning stub/template applications
- Finally, we are able to create a research applications repository that will continuously grow by accommodating machine learning augmented apps in areas such as medical diagnostics inference/prediction, financial predictive analytics, environment inferential models. The proposed scientific apps will be developed within the “Cloudifier” community and will fully utilize the scalable GPU-augmented machine learning template models;

III. PROPOSED ARCHITECTURE

The main points of our proposed architecture are the following:

- General execution framework based on REST services with HPC backend powered by GPU computing;
- Machine learning augmented applications generated within a provisioned Cloud infrastructure by automated software migration [4];
- Virtual desktop predictive analytics apps repository consisting of template-developed business applications and scientific experiments;

A. REST based massive parallel computation services

Based on the fact that one of our paper areas of innovation is the augmentation of shallow machine learning models with parallel processing techniques based on GPU processing, we are proposing a REST based service for parallel computation on GPU backend. At present time, based on our knowledge, the advancement of GPU based hardware infrastructures capable of delivering massive parallel computing for machine learning tasks at very low costs has only been exploited by the area of deep learning and deep neural networks in particular. Our proposal addresses this gap between the available computation resources for deep learning and the ones available for scientific shallow machine learning. In past years various papers have been arguing that shallow models can achieve similar results with deep models if sufficient data is provided, such as the famous quote defending the power of data by Google’s Research Director Peter Norvig: “We don’t have better algorithms. We just have more data.” [7]. In our proposed approach for advancing shallow machine learning by GPU parallel computing augmentation we will argue that shallow model’s pipelines can achieve comparable results with deep learning models for various applications. Although most of the state-of-the-art research in the area of machine learning is focused on deep learning models, the classic machine learning models with shallow hidden layer structure have two main characteristics that deep models lack: self-explain ability and algorithmic robustness for online and real-time environments.

Currently several programming approach options are available for programming GPU parallel computation engines both for scientific purposes and for engineering/production and probably the most widely used are OpenCL [6] and CUDA [7]. Basically, both mentioned models use the same technique with only slight differences. The two most fundamental aspect in regard to the computational approach consists in the two central entities consisting in *kernel* – that is the function containing the actual code to be run in parallel - and the *devices* – the actual computational devices that will run the *kernel* in parallel. Traditionally the *kernel* is a function described by the the following generic algorithm:

Algorithm 1 Kernel(*Data*, *BlockInformation*)

cID \leftarrow get coordinates tuple from *BlockInformation*
 Modify *Data*[cID] with hard coded *<operation>*

Return

Note that in above example *BlockInformation* defines the actual data segment that a particular *device* is processing within *Data* and *<operation>* denotes the hard-coded (compiled) operation that the *kernel* is computing for the particular block of *Data* it receives. For the particular case of matrix dot product we can have each *device* compute with its *kernel* the dot product of two vectors (the row vector of 1st matrix and the column vector of 2nd matrix).

$$A * B = A^T B = \sum a_i b_i \quad (1)$$

Finally, we have the following two code snippets for computing (1) in OpenCL and CUDA based *kernels*:

Code 1 OpenCL Kernel for Matrix-Matrix dot product

```
__kernel void OpenCLMatrixDotKernel(const int
A_ROWS,
const int BC_COLUMNS,
const int A_COLUMNS,
const __global float* A,
const __global float* B,
__global float* C)
{
    const int Row = get_global_id(0);
    const int Col = get_global_id(1);

    float acc = 0.0f;
    for (int c=0; c<A_COLUMNS; c++) {
        acc += A[c*A_ROWS + Row] *
B[Col*A_COLUMNS + c];
    }
    C[Col*A_ROWS + Row] = acc;
}
```

Code 2 CUDA Kernel for Matrix-Matrix dot product

```
__global__ void CUDAMatrixDotKernel(Matrix A,
Matrix B, Matrix C)
{
    float acc = 0;
    int row = blockIdx.y * blockDim.y +
threadIdx.y;
    int col = blockIdx.x * blockDim.x +
threadIdx.x;
    if(row > A.height || col > B.width)
        return;
    for (int e = 0; e < A.width; ++e)
        acc += A.elements[row * A.width +
e] *
            B.elements[e * B.width +
col];
    C.elements[row * C.width + col] = acc;
}
```

The implementation of the *kernel* code will be left for a lower level tensor computation framework, namely TensorFlow [6], a Open Source framework already established as a industry state-of-the-art standard. TensorFlow is a programming framework for expressing machine learning algorithms combined with a backend implementation for executing such algorithms in various heterogenous parallel computing environments. One of the most important features of TensorFlow is the automatic deployment of tensor graphs calculation on multiple GPUs, although without high-level resource allocation mechanics.

As previously mentioned, in the current paper we will present a REST [8] approach for cloud computing based provisioning of parallel computation services. This approach will enable the implementation and consumption of a web based service that will have the following generic algorithm based on FIFO allocation.

Algorithm 2 *ParallelProcessingREST(Request)*

For each Kernel, Data pair in Request

Kernel \Leftarrow get **high-end kernel** source code from *Request* (TensorFlow [6] code)

Data \Leftarrow get dataset from *Request* by value or by reference from external source

Available \Leftarrow get available devices from existing cluster based on adaptive **Algorithm 3**

If *Available* is valid then

Off-load *Kernel* and *Data* on *Available* devices

Get off-load *Result* and append to

Response

Else

Response \Leftarrow Overhead time

exceeds computing time

Return *Response*

The actual GPU execution strategy of the off-loaded kernels will be based on an adaptive algorithm that will take into consideration server-load, task complexity and GPU load/preparation overhead.

Algorithm 3 *ResourceAllocation(Kernel, Data)*

cFLOPS \Leftarrow evaluate needed FLOPS for *high-level Kernel* and *Data*

Resources \Leftarrow Compute needed devices based on *cFLOPS*

ComputeTime \Leftarrow Aproximate total execution time based on *Resources*, *cFLOPS*

WaitTime \Leftarrow Off-load overhead time + time until devices are available

If *ComputeTime* > *WaitTime*

Result \Leftarrow Queue position or actual allocated devices
else

Result \Leftarrow allocation invalidated

Return *Result*

Finally, the entire predictive analytics REST webservice and app repository is based on a simple algorithm (*Algorithm 4*) for servicing the following requests:

- Training new models based on post data or externally available data
- Prediction/inference based on trained models (via a model handle obtained after training)

Algorithm 4 *GetRequest(Handle, Data, Model)*

If *Handle* exists then

If *Data* denotes model status request:

Result \Leftarrow Model status (*Working*, *Trained*, *InferenceReady*)

Else

Prepare *Request* \Leftarrow *Handle*, *Data*, *Model*

Call *ParallelProcessingREST* (*Request*)

Update *Handle*

```

    Result  $\Leftarrow$  Handle
Else
  If
    Select high-level Kernel model  $\Leftarrow$  Model
    Call ParallelProcessingREST
    Generate Handle
    Result  $\Leftarrow$  Handle
Return Result

```

The actual messages that are passed between clients and the REST server are simple JSONs in the following format:

```

{
  HANDLE: ModelHandle, # not NULL if requesting
  results

  DATA: DataDict, # data dictionary containing
  predictor variables, targets (if available) with
  actual data or link to external data repository

  MODEL: ModelName # a model (specific or
  generic) to be trained (if model HANDLE is not
  yet available)
}

```

For a clearer understanding of the mechanics behind the proposed architecture we will present the full algorithm cycle for the case of requesting the training of a recommender system (*Use case algorithm 1*).

Use case algorithm 1 *Request(TransactionDataDict, "Recommender")*

```

Model  $\Leftarrow$  Determine the "Recommender" type of model
ParsedData  $\Leftarrow$  Parse the transactional data dict and
download data from external source if needed
HighLevelKernel  $\Leftarrow$  construct computational graph for
Model based on tensor computation/optimization
framework
TrainingJob  $\Leftarrow$  Allocate resources and request training of
Model based on ParsedData and HighLevelKernel
ModelTrainingHandle  $\Leftarrow$  obtain a handle of TrainingJob
Return ModelTrainingHandle

```

B. User Apps – migration and fast-development repository

The user created apps repository will consist of both automated generated apps and template-based apps. The generated apps will actually be the results of our automatic software migration engine [4], fully hosted within the execution environment proposed by our proposal.

The template-based apps purpose is to construct a continuously growing repository of online "skeleton" apps that have been designed for a particular horizontal purpose within the area of predictive analytics. The user will be able to configure a certain template or propose a new template for online

publication and crowd-testing. The proposed template apps cover the following types of predictive analytics models:

- Multi-variate machine learning regression model templates for product/sales recommendations
- Multi-variate machine learning classification model templates for customer churn prediction
- Unsupervised machine learning clustering model templates for customer segmentation

Finally, both the automated apps and the template based app creation features of our user apps repository will facilitate the fast-development and deployment of customer predictive analytics applications.

C. Business Apps and Research experiments

The business apps and research experiments repositories will host various predictive analytics online applications that will be able to run the proposed machine learning models within the proposed GPU processing backend. These repositories will contain applications that have been designed/developed using the API provided by the our cloud machine learning execution framework.

IV. CONCLUSIONS AND FUTURE DEVELOPMENTS

It is our objective to transform the experimental outcomes of the paper and the underlining architecture into a full community and crowd-sourced repository of both scientific and business-oriented machine learning experiments.

The main difference between the User Apps repository and the Business and Research Experiments repository is that in the first case the platform users are fully relying on our cloud architecture for the entire development lifecycle of the predictive analytics app. In the second case the platform users are externally developing the predictive analytics apps based on the resources provided by our cloud framework.

At present time, we are conducting experimentation activities in the following areas:

- Automatic legacy software migration and provisioning within proposed cloud infrastructure. A set of desktop-based ready-to-migrate applications has been constructed consisting of applications written in FoxPro, Visual Basic, Visual C++. For this particular set of applications, we are using our automated migration software based on [4] and the existing source code of the legacy where available. The source code is necessary for the actual validation of resulting automated migration inferential process.
- The second area of experimentation is the actual development and deployment of predictive analytics solutions for business applications. Currently customer behavior predictive analytics experiments are conducted within our proposed GPU-augmented machine learning framework based on our REST machine learning model

execution framework and TensorFlow [6] backend. Cross-validation and benchmarking of results is conducted using the well-known scientific and general machine learning frameworks such as:

- Python Sci-Kit Machine Learning package [11]
- R package “caret” (CRAN hosted) [12]

Conference on Machine Learning and Principles and Practices of Knowledge Discovery in Databases (2013), 2013.

- [12] M. Kuhn, "Caret package," *Journal of Statistical Software*, 2008.

V. BIBLIOGRAPHY

- [1] A. Agrawal, J. Gans and A. Goldfarb, "Managing the Machines," *HBR*, 2016.
- [2] Vinyals, "A Neural Conversational Model," in *ICML Deep Learning Workshop 2015*, 2015.
- [3] C. V, Q.-B. Nguyen and S. Pankanti, "Deep Learning Ensembles for Melanoma Recognition in Dermoscopy Images," *Journal of Research and Development*, 2016.
- [4] N. T. AI Damian, "Model Architecture for Automatic Translation and Migration of Legacy Applications to Cloud Computing Environments," in *CSCS*, Bucharest, 2017.
- [5] M. Bihis and S. Roychowdhury, "A generalized flow for multi-class and binary classification tasks: An Azure ML approach," in *Big Data (Big Data), 2015 IEEE International Conference on*, 2015.
- [6] M. Abadi, A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis, J. Dean, M. Devin, S. Ghemawat, I. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia and R. Jozefowicz, "TensorFlow: Large-Scale Machine Learning on Heterogeneous Distributed Systems," *Corenell University Library*; arXiv:1603.04467, 2016.
- [7] A. Halevy, P. Norvig and F. Pereira, "The Unreasonable Effectiveness of Data," *IEEE Intelligent Systems*, vol. 24, no. 2, pp. 8-12, 2009.
- [8] J. Stone, D. Gohara and G. Shi, "OpenCL: A Parallel Programming Standard for Heterogeneous Computing Systems," *Computing in Science & Engineering*, vol. 12, no. 3, 2010.
- [9] J. Ghorpade, J. Parande and M. Kulkarni, "GPGPU PROCESSING IN CUDA ARCHITECTURE," *Advanced Computing: An International Journal (ACIJ)*, vol. 3, 2012.
- [10] R. T. Fielding, "Architectural Styles and the Design of Network-based Software Architectures," *University of California*, Irvine, 2000.
- [11] L. B. (ILPS), G. Louppe, M. Blondel, F. P. (. S. -. I. d. France), A. Mueller, O. Grisel, V. Niculae, P. Prettenhofer, A. G. (. S. -. I. d. France, LTCI) and J. G. (. Saclay, "API design for machine learning software: experiences from the scikit-learn project," *European*